

Porter: Neural Networks Regression

Porter is India's Largest Marketplace for Intra-City Logistics. Leader in the country's \$40 billion intra-city logistics market, Porter strives to improve the lives of 1,50,000+ driver-partners by providing them with consistent earning & independence. Currently, the company has serviced 5+ million customers

Porter works with a wide range of restaurants for delivering their items directly to the people.

Problem statement

Porter has a number of delivery partners available for delivering the food, from various restaurants and wants to get an estimated delivery time that it can provide the customers on the basis of what they are ordering, from where and also the delivery partners.

This dataset has the required data to train a regression model that will do the delivery time estimation, based on all those features.

```
In [52]: import warnings
warnings.filterwarnings("ignore")
```

```
In [28]: import pandas as pd
import numpy as np
```

```
In [29]: df = pd.read_csv("dataset.csv")
df
```

Out[29]:

	market_id	created_at	actual_delivery_time	store_id	store_prio
0	1.0	2015-02-06 22:24:17	2015-02-06 23:27:16	df263d996281d984952c07998dc54358	
1	2.0	2015-02-10 21:49:25	2015-02-10 22:56:29	f0ade77b43923b38237db569b016ba25	
2	3.0	2015-01-22 20:39:28	2015-01-22 21:09:09	f0ade77b43923b38237db569b016ba25	
3	3.0	2015-02-03 21:21:45	2015-02-03 22:13:00	f0ade77b43923b38237db569b016ba25	
4	3.0	2015-02-15 02:40:36	2015-02-15 03:20:26	f0ade77b43923b38237db569b016ba25	
...
197423	1.0	2015-02-17 00:19:41	2015-02-17 01:24:48	a914ecef9c12ffdb9bede64bb703d877	
197424	1.0	2015-02-13 00:01:59	2015-02-13 00:58:22	a914ecef9c12ffdb9bede64bb703d877	
197425	1.0	2015-01-24 04:46:08	2015-01-24 05:36:16	a914ecef9c12ffdb9bede64bb703d877	
197426	1.0	2015-02-01 18:18:15	2015-02-01 19:23:22	c81e155d85dae5430a8cee6f2242e82c	
197427	1.0	2015-02-08 19:24:33	2015-02-08 20:01:41	c81e155d85dae5430a8cee6f2242e82c	

197428 rows × 14 columns



In [30]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 197428 entries, 0 to 197427
Data columns (total 14 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   market_id                            196441 non-null  float64
1   created_at                           197428 non-null  object
2   actual_delivery_time                 197421 non-null  object
3   store_id                             197428 non-null  object
4   store_primary_category               192668 non-null  object
5   order_protocol                       196433 non-null  float64
6   total_items                          197428 non-null  int64
7   subtotal                            197428 non-null  int64
8   num_distinct_items                  197428 non-null  int64
9   min_item_price                      197428 non-null  int64
10  max_item_price                      197428 non-null  int64
11  total_onshift_partners               181166 non-null  float64
12  total_busy_partners                  181166 non-null  float64
13  total_outstanding_orders             181166 non-null  float64
dtypes: float64(5), int64(5), object(4)
memory usage: 21.1+ MB
```

```
In [31]: #Converting the date-time to datatype
df['actual_delivery_time'] = pd.to_datetime(df['actual_delivery_time'])
df['created_at'] = pd.to_datetime(df['created_at'])
```

```
In [32]: #Delivery time in minutes
df['delivery_time'] = (df['actual_delivery_time'] - df['created_at']).dt.total_seconds() / 60
df['hour'] = df['created_at'].dt.hour
df['day'] = df['created_at'].dt.dayofweek

df.market_id = df.market_id.astype('category')
df.order_protocol = df.order_protocol.astype('category')
```

```
In [33]: df.isna().sum()
```

```
Out[33]: market_id                987
created_at                0
actual_delivery_time       7
store_id                  0
store_primary_category    4760
order_protocol             995
total_items                0
subtotal                  0
num_distinct_items        0
min_item_price             0
max_item_price             0
total_onshift_partners    16262
total_busy_partners       16262
total_outstanding_orders  16262
delivery_time              7
hour                       0
day                        0
dtype: int64
```

Deleting all the rows with null values.

```
In [34]: df.dropna(inplace=True)
df.shape
```

```
Out[34]: (176248, 17)
```

```
In [35]: #Converting float columns to int datatype
float_cols = ['total_onshift_partners', 'total_busy_partners', 'total_outstanding_order']
df[float_cols] = df[float_cols].astype('int')
```

Deleting columns which are not required

```
In [36]: df.drop(columns=['store_id', 'created_at', 'actual_delivery_time'], inplace=True)
```

```
In [37]: df.describe()
```

```
Out[37]:
```

	market_id	order_protocol	total_items	subtotal	num_distinct_items	min_item_
count	176248.000000	176248.000000	176248.000000	176248.000000	176248.000000	176248.000000
mean	2.743747	2.911687	3.204592	2696.498939	2.674589	684.91
std	1.330911	1.512920	2.673899	1828.922584	1.625558	519.91
min	1.000000	1.000000	1.000000	0.000000	1.000000	-86.00
25%	2.000000	1.000000	2.000000	1408.000000	1.000000	299.00
50%	2.000000	3.000000	3.000000	2221.000000	2.000000	595.00
75%	4.000000	4.000000	4.000000	3407.000000	3.000000	942.00
max	6.000000	7.000000	411.000000	26800.000000	20.000000	14700.00

```
In [39]: #removing negative values
df = df[~((df[['total_onshift_partners', 'total_busy_partners', 'total_outstanding_order'] < 0)))]
```

```
In [41]: df.describe()
```

```
Out[41]:
```

	market_id	order_protocol	total_items	subtotal	num_distinct_items	min_item_
count	176157.000000	176157.000000	176157.000000	176157.000000	176157.000000	176157.000000
mean	2.744109	2.911789	3.204079	2696.578484	2.674580	684.91
std	1.330883	1.512871	2.671920	1829.049610	1.625497	519.81
min	1.000000	1.000000	1.000000	0.000000	1.000000	0.00
25%	2.000000	1.000000	2.000000	1408.000000	1.000000	299.00
50%	2.000000	3.000000	3.000000	2221.000000	2.000000	595.00
75%	4.000000	4.000000	4.000000	3408.000000	3.000000	942.00
max	6.000000	7.000000	411.000000	26800.000000	20.000000	14700.00

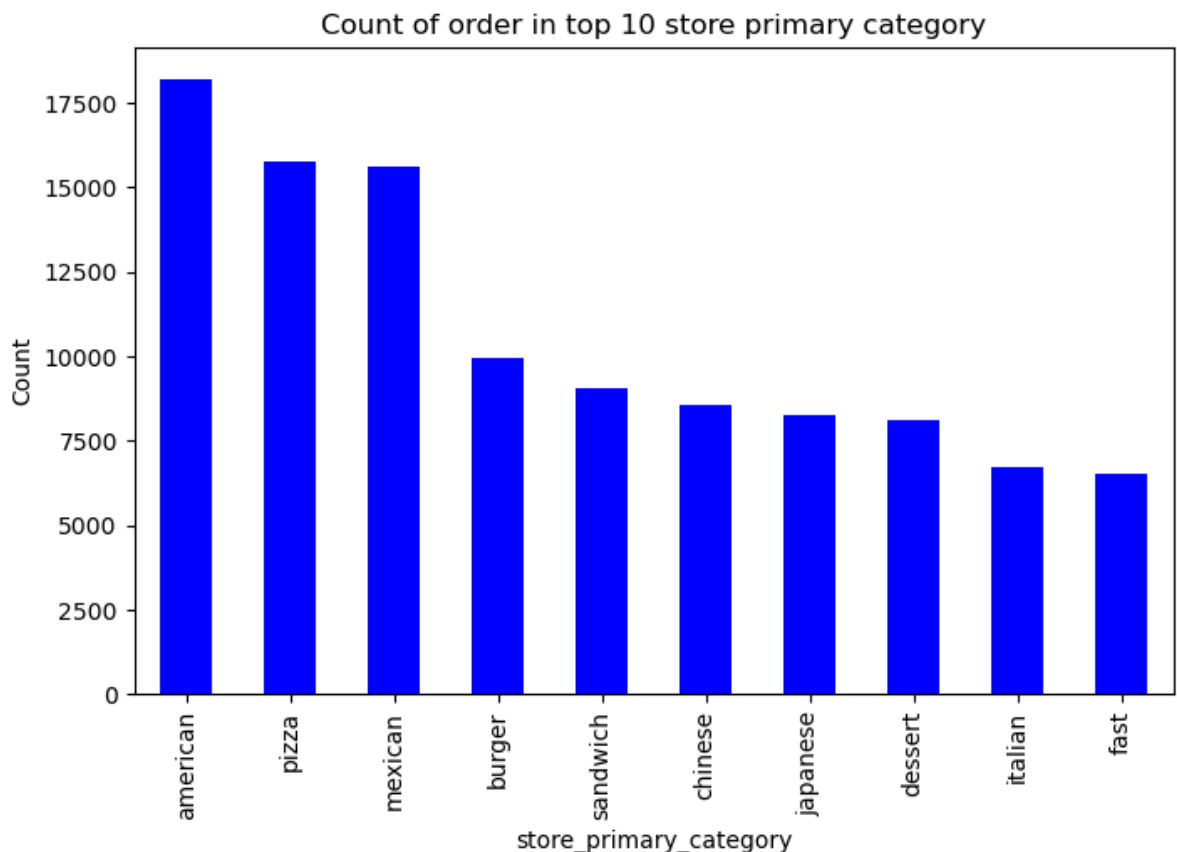
```
In [40]: df.isna().sum()
```

```
Out[40]: market_id      0
store_primary_category  0
order_protocol          0
total_items            0
subtotal              0
num_distinct_items     0
min_item_price         0
max_item_price         0
total_onshift_partners  0
total_busy_partners    0
total_outstanding_orders 0
delivery_time          0
hour                  0
day                   0
dtype: int64
```

Univariate Bivariate and Multivariate Analysis

```
In [43]: fig = plt.figure(figsize=(8,5))

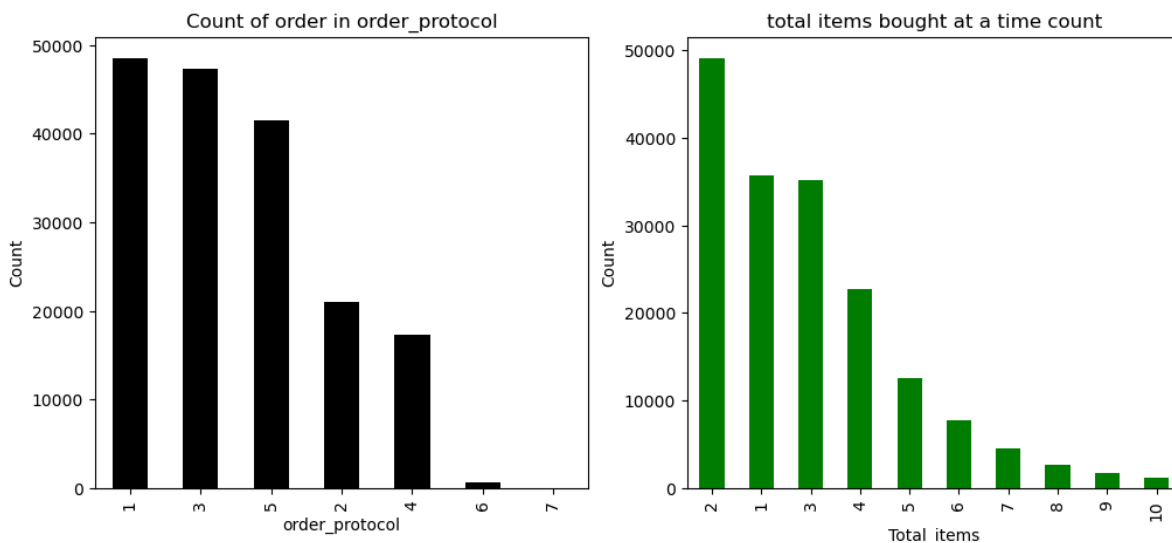
df['store_primary_category'].value_counts().head(10).plot(kind='bar',color='blue')
plt.xlabel('store_primary_category')
plt.ylabel('Count')
plt.title('Count of order in top 10 store primary category')
plt.show()
```



```
In [46]: fig = plt.figure(figsize=(12,5))
plt.subplot(1,2,1)
df['order_protocol'].value_counts().plot(kind='bar',color='black')
plt.xlabel('order_protocol')
plt.ylabel('Count')
plt.title('Count of order in order_protocol')

plt.subplot(1,2,2)
df['total_items'].value_counts().head(10).plot(kind='bar',color='green')
```

```
plt.xlabel('Total_items')
plt.ylabel('Count')
plt.title('total items bought at a time count')
plt.show()
```




```
In [49]: # Distribution of Categorical Features

fig, ax = plt.subplots(2,1, figsize=(8,8))

sns.countplot(data=df, x='hour', palette='Blues', ax=ax[0])
ax[0].set_title('Order Hour')

sns.countplot(data=df, x='day', palette='Blues', ax=ax[1])
ax[1].set_title('Order Day')

plt.tight_layout()
plt.show()
```



```
In [50]: fig, ax = plt.subplots(2,4, figsize=(12,8))

df.boxplot(column='total_items', ax=ax[0,0])
ax[0,0].set_title('Total Items')

df.boxplot(column='subtotal', ax=ax[0,1])
ax[0,1].set_title('Subtotal')

df.boxplot(column='min_item_price', ax=ax[0,2])
ax[0,2].set_title('Min Price')

df.boxplot(column='max_item_price', ax=ax[0,3])
ax[0,3].set_title('Max Price')

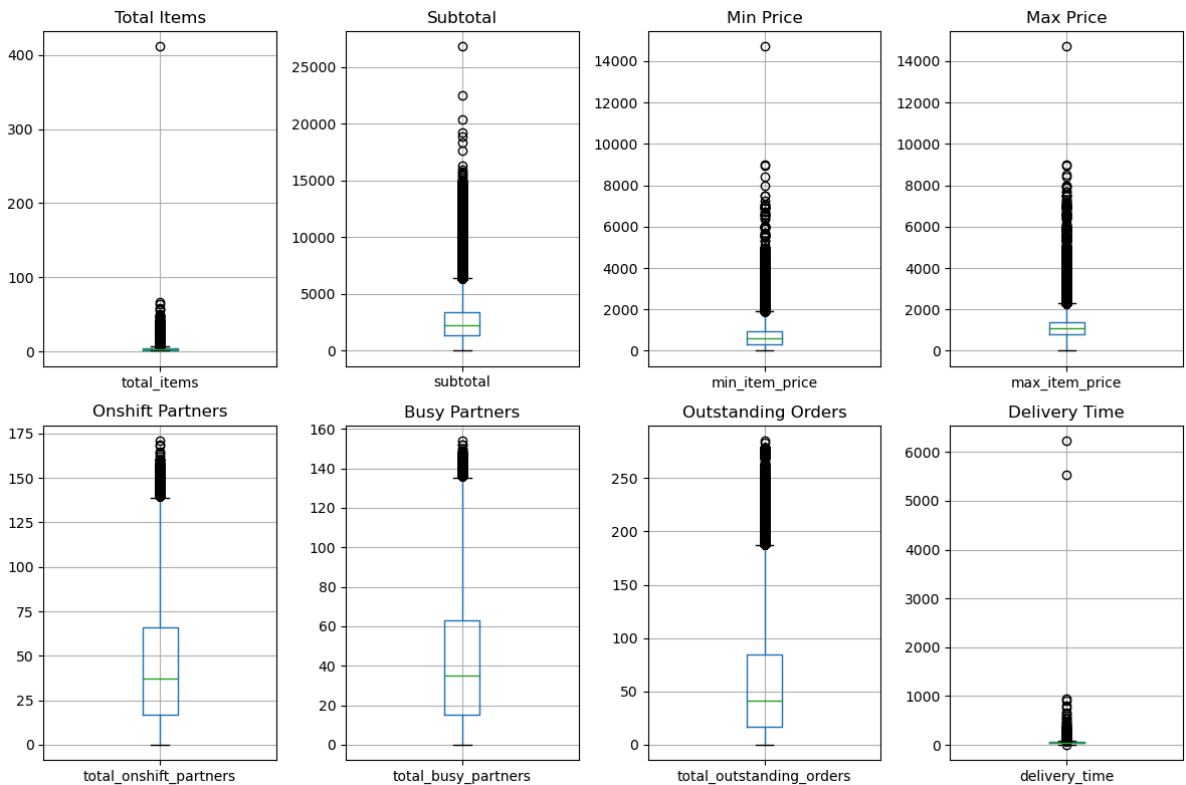
df.boxplot(column='total_onshift_partners', ax=ax[1,0])
ax[1,0].set_title('Onshift Partners')

df.boxplot(column='total_busy_partners', ax=ax[1,1])
ax[1,1].set_title('Busy Partners')

df.boxplot(column='total_outstanding_orders', ax=ax[1,2])
ax[1,2].set_title('Outstanding Orders')

df.boxplot(column='delivery_time', ax=ax[1,3])
ax[1,3].set_title('Delivery Time')
```

```
plt.tight_layout()
plt.show()
```



Outlier Treatment

```
In [53]: # Remove outliers through IQR method
outlier_cols = ['total_items', 'subtotal', 'min_item_price',
                'max_item_price', 'delivery_time', 'total_onshift_partners',
                'total_busy_partners', 'total_outstanding_orders']

for col in outlier_cols:
    # Calculate Q1 and Q3
    q1 = np.percentile(df[col], 25)
    q3 = np.percentile(df[col], 75)

    # Calculate IQR
    iqr = q3 - q1


    # Define lower and upper bounds
    lower_bound = q1 - (1.5 * iqr)
    upper_bound = q3 + (1.5 * iqr)

    # Remove outliers
    df = df.loc[~((df[col] < lower_bound) | (df[col] > upper_bound))]

df.shape
```


Out[53]: (142433, 14)

```
In [55]: plt.figure(figsize=(10,5))
sns.histplot(x='subtotal', data=df, kde=True, edgecolor='red')
plt.show()
```

```
In [56]: import seaborn as sns
import matplotlib.pyplot as plt
sns.heatmap(df.corr())
```

Out[56]: <Axes: >



1. Most orders were for restaurants in the American, Pizza, Mexican, and Burger categories, in that order.

2. The majority of orders followed protocol 1, with protocols 3 and 5 following. Protocol 7 saw very few orders.
3. Most orders were for 2 items at a time, followed by 1, 3, 4, and so on.
4. No orders were placed between 9 and 13 hours, and very few orders were placed at 6, 7, 8, 14, 15, and 16 hours. It appears that the hours between 7 and 16 don't see many orders.
5. Weekends had a higher order count.
6. There was almost no difference between the count of orders based on the time of delivery and the time of order creation, suggesting that active hours are between 0-5 and 17-23.
7. No clear trend was observed based on the hour of the day or day of the week, aside from some off-peak times.
8. No correlation could be found between the total time taken to deliver an item and other numerical features.
9. There is a strong correlation between the following feature pairs: (total_onshift_partners, total_busy_partners, total_outstanding_orders), (hour_created, hour_delivered), and (day_created, day_delivered).

Data Preparation

```
In [57]: X = df.drop(columns=['delivery_time'])
y = df['delivery_time']

print(X.shape, y.shape)

(142433, 13) (142433,)
```

```
In [61]: from sklearn.preprocessing import OneHotEncoder

cat_cols = ['market_id', 'order_protocol', 'hour', 'day']
enc = OneHotEncoder(handle_unknown='ignore')
X_encoded = pd.DataFrame(enc.fit_transform(X[cat_cols]).toarray(), columns=enc.get_feature_names_out(cat_cols))
X_num = X.drop(cat_cols, axis=1)
X = pd.concat([X_num, X_encoded], axis=1)
```

```
In [63]: #split the data for training, validation and test

from sklearn.model_selection import train_test_split

X_train_val, X_test, y_train_val, y_test = train_test_split(X,y,test_size=0.1, random_state=42)
X_train, X_val, y_train, y_val = train_test_split(X_train_val, y_train_val, test_size=0.5, random_state=42)

print('Train: ', X_train.shape, y_train.shape)
print('Val: ', X_val.shape, y_val.shape)
print('Test: ', X_test.shape, y_test.shape)

Train: (115370, 47) (115370,)
Val: (12819, 47) (12819,)
Test: (14244, 47) (14244,)
```

Encoding

```
In [64]: from category_encoders import TargetEncoder

encoder = TargetEncoder(cols=['store_primary_category'])

X_train = encoder.fit_transform(X_train, y_train)
X_val = encoder.transform(X_val)
X_test = encoder.transform(X_test)
```

```
In [65]: from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_train = pd.DataFrame(X_train_scaled, columns=X_train.columns)

X_val_scaled = scaler.transform(X_val)
X_val=pd.DataFrame(X_val_scaled, columns=X_val.columns)

X_test_scaled = scaler.transform(X_test)
X_test=pd.DataFrame(X_test_scaled, columns=X_test.columns)
```

Neural Network

```
In [69]: import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout
from tensorflow.keras.optimizers import Adam

# Model definition
model = Sequential()
model.add(Dense(64, kernel_initializer='normal', activation='relu', input_shape=(X_
model.add(Dense(512, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(256, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(1, activation='linear'))


# Model compilation
model.compile(loss='mse', optimizer=Adam(), metrics=['mae'])

# Model training
history = model.fit(X_train, y_train, epochs=30, batch_size=512, verbose=1, validat
```

Epoch 1/30

226/226  6s 17ms/step - loss: 658.9433 - mae: 19.4586 - val_loss: 167.9703 - val_mae: 10.2440

Epoch 2/30

226/226  3s 14ms/step - loss: 168.4960 - mae: 10.2641 - val_loss: 161.3014 - val_mae: 10.1014

Epoch 3/30

226/226  3s 15ms/step - loss: 163.6679 - mae: 10.1344 - val_loss: 160.3808 - val_mae: 10.0292

Epoch 4/30

226/226  3s 14ms/step - loss: 161.0559 - mae: 10.0285 - val_loss: 159.3210 - val_mae: 9.9441

Epoch 5/30

226/226  3s 14ms/step - loss: 160.6059 - mae: 10.0424 - val_loss: 158.8936 - val_mae: 9.8771

Epoch 6/30

226/226  3s 14ms/step - loss: 159.1536 - mae: 9.9747 - val_loss: 159.3620 - val_mae: 9.8863


Epoch 7/30

226/226  3s 14ms/step - loss: 160.3427 - mae: 9.9939 - val_loss: 157.6386 - val_mae: 9.9454

Epoch 8/30

226/226  3s 14ms/step - loss: 160.7506 - mae: 10.0244 - val_loss: 158.6249 - val_mae: 10.1010

Epoch 9/30

226/226  3s 15ms/step - loss: 158.5397 - mae: 9.9619 - val_loss: 157.6664 - val_mae: 9.8620


Epoch 10/30

226/226  3s 14ms/step - loss: 157.7243 - mae: 9.9233 - val_loss: 159.2165 - val_mae: 9.8683

Epoch 11/30

226/226  5s 15ms/step - loss: 158.1557 - mae: 9.9198 - val_loss: 157.7001 - val_mae: 9.9731

Epoch 12/30

226/226  3s 15ms/step - loss: 159.4295 - mae: 9.9711 - val_loss: 156.7899 - val_mae: 9.8632


Epoch 13/30

226/226  3s 14ms/step - loss: 156.5410 - mae: 9.8854 - val_loss: 159.6482 - val_mae: 9.8230

Epoch 14/30

226/226  3s 14ms/step - loss: 156.9104 - mae: 9.8990 - val_loss: 157.5926 - val_mae: 9.8704

Epoch 15/30

226/226  3s 14ms/step - loss: 157.3654 - mae: 9.8923 - val_loss: 157.6852 - val_mae: 9.8252

Epoch 16/30

226/226  3s 14ms/step - loss: 155.3760 - mae: 9.8367 - val_loss: 157.7569 - val_mae: 9.8277

Epoch 17/30

226/226  3s 14ms/step - loss: 156.4272 - mae: 9.8699 - val_loss: 156.0047 - val_mae: 9.9212

Epoch 18/30

226/226  3s 15ms/step - loss: 156.7515 - mae: 9.8912 - val_loss: 158.3524 - val_mae: 9.7995

Epoch 19/30

226/226  3s 15ms/step - loss: 155.1786 - mae: 9.8318 - val_loss: 157.7239 - val_mae: 9.7945

Epoch 20/30

226/226  3s 15ms/step - loss: 156.1209 - mae: 9.8580 - val_loss: 157.0641 - val_mae: 9.8153

Epoch 21/30

226/226  3s 14ms/step - loss: 156.5166 - mae: 9.8723 - val_loss: 156.4756 - val_mae: 9.9386

Epoch 22/30

226/226 ————— 3s 14ms/step - loss: 155.9741 - mae: 9.8615 - val_loss: 156.2014 - val_mae: 9.8319
Epoch 23/30

226/226 ————— 3s 14ms/step - loss: 155.3581 - mae: 9.8586 - val_loss: 156.1990 - val_mae: 9.8898
Epoch 24/30

226/226 ————— 3s 14ms/step - loss: 155.4746 - mae: 9.8463 - val_loss: 156.3251 - val_mae: 9.8419
Epoch 25/30

226/226 ————— 3s 14ms/step - loss: 155.5638 - mae: 9.8596 - val_loss: 157.2416 - val_mae: 10.0491
Epoch 26/30

226/226 ————— 3s 14ms/step - loss: 155.8516 - mae: 9.8623 - val_loss: 157.8698 - val_mae: 9.8325
Epoch 27/30


226/226 ————— 3s 15ms/step - loss: 154.6075 - mae: 9.7978 - val_loss: 155.7472 - val_mae: 9.8931
Epoch 28/30

226/226 ————— 3s 14ms/step - loss: 154.1745 - mae: 9.8131 - val_loss: 156.1086 - val_mae: 9.9201
Epoch 29/30

226/226 ————— 3s 14ms/step - loss: 155.7748 - mae: 9.8496 - val_loss: 155.9931 - val_mae: 9.8565
Epoch 30/30

226/226 ————— 3s 14ms/step - loss: 154.5389 - mae: 9.8037 - val_loss: 155.8954 - val_mae: 9.9424

```
In [70]: plt.plot(history.history['loss'], label='training data')
plt.plot(history.history['val_loss'], label='validation data')
plt.title('Training and Validation Loss across Epochs')
plt.ylabel('MSE value')
plt.xlabel('epochs')
plt.legend(loc='upper right')
plt.show()
```



In [74]: *#Evaluating the Model*

```
val_loss, val_mae = model.evaluate(X_val, y_val, verbose=1)
print(f'Validation Loss: {val_loss}, Validation MAE: {val_mae}')
```

401/401 ————— 1s 2ms/step - loss: 156.0142 - mae: 9.9293
 Validation Loss: 155.89535522460938, Validation MAE: 9.94243335723877

- the model is not overfitting and is performing consistently on both the training and validation datasets.
- the model generalizes well, with similar prediction errors on unseen data

In [76]: **from** sklearn.metrics **import** mean_squared_error, mean_absolute_error

```
y_pred = model.predict(X_val)
```

```
# Mean Squared Error (MSE)
```

```
mse = mean_squared_error(y_val, y_pred)
print(f'Mean Squared Error (MSE): {mse}')
```

```
# Root Mean Squared Error (RMSE)
```

```
rmse = np.sqrt(mse)
print(f'Root Mean Squared Error (RMSE): {rmse}')
```

```
# Mean Absolute Error (MAE)
```


```
mae = mean_absolute_error(y_val, y_pred)
print(f'Mean Absolute Error (MAE): {mae}')
```

401/401 ————— 1s 2ms/step
 Mean Squared Error (MSE): 155.89536218754068
 Root Mean Squared Error (RMSE): 12.48580642920355
 Mean Absolute Error (MAE): 9.942436836844385

In [75]: *#Predicting and Analyzing Results*

```
plt.scatter(y_val, y_pred)
plt.xlabel('Actual Values')
plt.ylabel('Predicted Values')
plt.title('Actual vs Predicted')
plt.show()
```

401/401 ————— 1s 2ms/step



Recommendations:

1. **Incentives for Night Shifts:** Given the high delivery demand between midnight and 5 AM and the limited number of delivery personnel available at night, Porter should offer incentives and bonuses for night shifts. This can enhance customer satisfaction by meeting the substantial late-night demand.
1. **Weekend Staffing:** Since weekends see larger order volumes, Porter should consider hiring employees specifically for weekend shifts. This approach allows scaling up operations without the need for permanent staff.
2. **Strategic Placement of Delivery Personnel:** To improve delivery times and customer satisfaction, Porter should position delivery personnel near popular restaurants, such as those serving American cuisine and pizza, which receive the majority of orders.
3. **Promote Underutilized Portals:** Portal 1 is receiving most of the orders, while Portal 7 remains underutilized. Porter should increase awareness and encourage customers to use Portal 7 to balance order distribution across platforms.

Leading Questions:

1. **Defining the problem statements and where can this and modifications of this be used?**

Problem Statement: Porter aims to develop a regression model to estimate delivery times for food orders based on various factors such as the type of order, the restaurant, and the

delivery partner. This involves predicting how long it will take for an order to be delivered from the restaurant to the customer.

Use Cases and Modifications:

- **Customer Communication:** Accurate delivery time predictions can improve customer satisfaction by setting realistic expectations.
- **Operational Efficiency:** Helps optimize delivery schedules and resource allocation, potentially reducing costs and enhancing efficiency.
- **Dynamic Pricing:** Delivery times can be factored into pricing models, adjusting costs based on expected delivery duration and urgency.

2. List 3 functions the pandas datetime provides with one line explanation.

1. **pd.to_datetime():** Converts a scalar, array-like, or Series into a pandas datetime object.
2. **pd.DatetimeIndex():** Creates a DatetimeIndex object from an array-like or iterable object of date-time strings.
3. **pd.date_range():** Generates a fixed frequency date range, which is useful for creating sequences of dates.

3. Short note on datetime, timedelta, time span (period)

Datetime: Represents a specific point in time, including both date and time components. In pandas, it is used for time-series data, enabling easy manipulation and analysis of temporal data.

Timedelta: Represents a duration or difference between two datetime values. It is used to perform operations like adding or subtracting time intervals from datetime objects.

Time Span (Period): Represents a specific length of time or a range between two dates. In pandas, the Period class can be used to represent and work with periods of time, such as months or years, allowing for easy manipulation of time periods.

4. Why do we need to check for outliers in our data?

- Due to outliers, scaling techniques such as standardisation or MinMaxScaler are impacted and result in wrong representation of the data for training
- Outliers can significantly impact descriptive statistics of the data like mean, standard deviation and correlation. This can cause the data to be misrepresented
- Outliers disproportionately influence the weights and biases while model training leading to overfitting. Outliers in training data make it difficult for the model to generalise well on test data.

5. Name 3 outlier removal methods?

Z-Score Method, IQR (Interquartile Range) Method and standard deviation method.

6. What classical machine learning methods can we use for this problem?

For estimating delivery times, classical machine learning methods like **Linear Regression** provide a straightforward approach, while **Decision Trees** and **Random Forests** handle complex, non-linear relationships effectively. **Gradient Boosting Machines (GBM)** and

Support Vector Regression (SVR) offer advanced techniques for capturing intricate patterns and improving prediction accuracy.

7. Why is scaling required for neural networks?

Scaling is required for neural networks to ensure that input features are on a similar scale, which helps in achieving faster convergence during training. It also prevents issues with gradient descent where features with larger magnitudes can dominate, leading to inefficient learning or instability in the model's training process.

8. Briefly explain your choice of optimizer.

Adam (Adaptive Moment Estimation): Combines the advantages of two other optimizers, Momentum and RMSprop. It adjusts the learning rate for each parameter individually based on estimates of first and second moments of the gradients, which helps in converging faster and handling sparse gradients effectively. It is widely used due to its adaptive learning rate and robustness to hyperparameter settings.

9. Which activation function did you use and why?

ReLU (Rectified Linear Unit) was used as the activation function because it introduces non-linearity into the model while being computationally efficient. ReLU helps in mitigating the vanishing gradient problem by allowing gradients to flow through the network more effectively during training, which accelerates convergence and improves performance.

10. Why does a neural network perform well on a large dataset?

A neural network performs well on a large dataset because:

1. **Better Generalization:** Large datasets provide diverse examples, helping the model learn more generalized patterns and reduce overfitting, leading to improved performance on unseen data.
2. **Complexity Handling:** Neural networks have numerous parameters and layers, and large datasets help in training these complex models effectively by providing sufficient data to capture intricate relationships and patterns.